

SITE-PROBE ACCURACY IS NOT A HARMONIZATION SCORE: AN AUDIT OF EVALUATION PRACTICE FOR MULTI-SITE MRI

Muhammad Muqsit Islam^{1,2,*} Gloria García Cuenco^{1,2,†}
José M. Martínez Sánchez¹ Pierrick Coupé²

¹ Universidad Autónoma de Madrid, Spain

² University of Bordeaux, France

ABSTRACT

Image-level MRI harmonization is often scored by how poorly a frozen site classifier recognizes the acquisition site of harmonized images. We audit this practice on raw ABIDE I T1-weighted MRI (1,112 subjects, 17 sites), using ten outputs from seven harmonization families as test cases and paired subject-level intervals. (i) Verdicts depend on how the probe is trained: five probes trained with the benchmark’s recipe gave the same CycleGAN outputs site balanced accuracies of 0.13–0.59; five converged probes agreed (0.44–0.54) but found more site information in eight of ten outputs, e.g. 0.70 versus 0.21 for aggressive StarGAN. (ii) The lowest accuracy, at chance, came from the most destructive output. (iii) Diffusion translation changed intensities least but barely moved them toward the target site. (iv) Probes trained on harmonized outputs recovered the source site from whole heads, skull-stripped brains and brain intensity histograms alike. We propose a reporting checklist.

Index Terms— MRI harmonization, evaluation, site classification, domain shift, reproducibility

1. INTRODUCTION

Acquisition differences between sites confound pooled MRI analyses and degrade models on unseen scanners [1, 2]. Image-level harmonization translates scans toward a common appearance. Travelling subjects [3] or paired protocols [4] allow direct checks but cover few subjects and scanners, so methods applied to large retrospective cohorts are judged indirectly [5]. A common proxy is a *site probe*: a classifier trained to recognize the site on unharmonized images and applied, frozen, to harmonized ones [2, 6]. Lower accuracy is read as removal of site information and is paired with similarity to the input, such as PSNR, and with intensity-distribution distances.

Each proxy measures less than it appears to. A frozen probe reports what one classifier recovers; it can also fail

because outputs are unfamiliar, damaged or noisy. Removing an attribute until a fixed adversary fails does not stop a newly trained classifier from recovering it [7], and for harmonized features site classifiers retrained after ComBat still predicted site above chance [8]. Site is also decodable at about 0.9 balanced accuracy from frozen, even randomly initialized, encoders, so the signal lies partly in low-level image statistics [9]. Similarity to the input rewards doing nothing, and a small input-to-output distance says nothing about the target site.

We audit these practices on outputs of seven harmonization families. Contributions: (1) a subject-level protocol with paired uncertainty, valid probe controls, separate reporting of target-site subjects, and target alignment measured apart from change; (2) evidence that probe verdicts depend on how the probe was trained, that the lowest site accuracy can signal destruction, and that small change does not imply alignment; (3) a test of removed versus hidden site information using probes trained on harmonized images, with geometry-only, intensity-only and brain-only controls; and (4) a reporting checklist.¹ We neither propose nor rank harmonizers.

2. DATA AND AUDITED OUTPUTS

Cohort. Raw ABIDE I T1-weighted scans [10] of 1,112 subjects from 17 sites are split by subject, stratified by site, into 890 training, 113 validation and 109 test subjects. Volumes are reoriented, clipped to $[p_1, p_{99}]$ and scaled to $[0, 1]$; one axial slice per subject is chosen by foreground area, cropped to a dilated head mask and resized to 256×256 . NYU, the largest site, is the target; 19 test subjects are from NYU and 90 from 16 source sites. Exact ties made slice choice host-dependent, so slices are frozen and verified for every output. The label is coarser than the acquisition: ABIDE I releases UM, UCLA and Leuven as two sub-samples each, merged by our labels, and 11 of 17 sites hold more than one voxel size or matrix (UCLA includes 10 scans at $1.5 \times 1.5 \times 4$ mm); no per-subject scanner identifier is released.

*Corresponding author

† Equal contribution

¹Code, the evaluation protocol with all amendments, and per-subject results: <https://github.com/muqsitamir/HarmonIt>.

Audited outputs. Earlier benchmark outputs were evaluated without retraining: NeuroCombat [11, 1] per pixel, fit on the test cohort itself (transductive); histogram matching of foreground quantiles to a pooled 17-site training reference [12]; CycleGAN [13]; StarGAN [14], aggressive and conservative; DLEST [15] after 1,000 and 1,500 steps; conditional diffusion image-to-image translation [16, 17] (strength 0.35, 50 DDIM steps); and adapted volumetric HCLD [18]. Learned methods were trained on the training split. Test subjects were inspected in the earlier benchmark, so this is a retrospective reanalysis; each audit experiment was written into a protocol before its outcome was seen, recording what had been observed.

3. AUDIT PROTOCOL

Change and target alignment. PSNR on a fixed $[0, 1]$ scale and normalized cross-correlation (XCorr) compare output with input on whole images; generators return NYU inputs almost unchanged (PSNR infinite or above 150 dB), so we summarize the 90 source subjects. *Change* is the per-subject Wasserstein-1 distance W between input and output intensities. *Target alignment* is the Wasserstein-1 distance and KL divergence between foreground intensities (defined on the raw slice) and a reference from the 147 NYU training subjects; ΔW_{NYU} is the output’s distance minus the input’s.

Site probes. All probes are single-channel ResNet-18 networks [19] predicting 17 sites, trained with AdamW (learning rate 3×10^{-4} , batch 64, 500 steps) and random affine augmentation, each selected on validation BA from its own training source and evaluated once. We train 44 of them in five families:

Frozen (1): the earlier benchmark’s own checkpoint.

Benchmark recipe (5 seeds): that training repeated on raw volumes. Does the seed change the verdict?

Converged (5 seeds): four times the steps with cosine decay, tested at the final epoch. Does a converged probe agree?

Shuffled labels (4): labels permuted between subjects. Is the probe reading images, not labels?

Slice probes (24, plus 6 restricted-input networks and 8 logistic fits): one fixed slice per subject, from raw images or from one method’s outputs. Removed or hidden? Restricted variants see only head silhouettes or brain masks (geometry), only foreground intensity histograms (intensity), or slices masked by the raw volume’s HD-BET brain mask [20].

The three methods behind the slice probes were chosen in advance to span intensity, adversarial and diffusion families. The converged, silhouette, histogram and brain-only experiments were added after earlier results and before their own, as the protocol records.

Statistics. Source BA is macro-averaged over the 16




Fig. 1. (a) Source site BA of each output, ordered by PSNR; retrained probes show the mean and range over five seeds. (b, c) Probes trained on raw slices (grey) or on a method’s outputs (aqua), tested on that method’s outputs (diffusion: second draw), on whole heads and on brain-only slices: image probes (circles, mean and range over 3 seeds) and intensity-histogram probes (squares, 95% interval). Black bar: raw-trained probe on raw test slices. Grey: image probes given only head silhouettes or brain masks.

source sites (chance 1/16). Intervals are 95% percentiles of 2,000 bootstrap replicates resampling subjects within site, shared across methods and probes (paired) and conditional on fixed checkpoints.

4. RESULTS

Probe validity. On raw source images the frozen probe reached BA 0.95 [0.90, 1.00], benchmark-recipe probes 0.81–0.98 and converged probes 1.00; shuffled-label controls scored 0.03. Probes learn site-related structure, not necessarily scanner appearance: head silhouettes alone gave 0.73–0.86.

Verdicts depend on probe training. Benchmark-recipe probes assigned very different site BA to the same outputs

Table 1. Source-subject results ($n = 90$). PSNR and XCorr compare output with input; ΔW_{NYU} : change in Wasserstein-1 distance to the NYU reference (negative is closer). Site BA: frozen probe with 95% interval; five benchmark-recipe and five converged probes as mean [min, max]. Chance is 0.06. [†]Fit on the test cohort. [‡]Pooled 17-site reference.

Method	Change and target alignment			Source-site balanced accuracy of probes trained on raw images		
	PSNR	XCorr	ΔW_{NYU}	Frozen probe	Benchmark recipe	Converged recipe
Raw input	—	—	0	0.95	0.90 [0.81, 0.98]	1.00 [1.00, 1.00]
NeuroCombat [†]	20.0	0.940	+0.003	0.27 [0.21, 0.32]	0.33 [0.26, 0.37]	0.34 [0.28, 0.42]
Histogram matching [‡]	23.0	0.982	-0.053	0.44 [0.37, 0.52]	0.44 [0.27, 0.59]	0.61 [0.55, 0.68]
CycleGAN	21.3	0.977	-0.036	0.31 [0.25, 0.37]	0.28 [0.13, 0.59]	0.49 [0.44, 0.54]
StarGAN (aggr.)	16.0	0.896	-0.019	0.21 [0.15, 0.27]	0.35 [0.18, 0.54]	0.70 [0.65, 0.81]
StarGAN (cons.)	26.9	0.991	-0.021	0.83 [0.79, 0.86]	0.78 [0.66, 0.89]	0.98 [0.94, 0.99]
DLEST 1500	26.6	0.993	-0.020	0.67 [0.61, 0.72]	0.73 [0.65, 0.82]	0.96 [0.95, 0.97]
DLEST 1000	28.4	0.993	-0.009	0.71 [0.65, 0.77]	0.84 [0.76, 0.89]	0.99 [0.96, 1.00]
Diffusion img2img	21.7	0.948	-0.001	0.37 [0.32, 0.42]	0.42 [0.24, 0.63]	0.57 [0.52, 0.59]
Adapted HCLD	13.9	0.706	+0.108	0.07 [0.06, 0.08]	0.08 [0.06, 0.11]	0.08 [0.06, 0.10]




Fig. 2. Source test subject with median PSNR over all outputs; insets are a $2\times$ zoom of the boxed region. Bottom: output minus input, with this subject's PSNR. HCLD loses anatomy; diffusion adds texture that differs between draws.

(Table 1, Fig. 1a): CycleGAN 0.13–0.59 and aggressive StarGAN 0.18–0.54, with Kendall τ 0.51–0.87 between method orderings. Their validation BA fluctuated between epochs (e.g. 0.88, 0.50, 0.87), and raw accuracy did not reveal it: the probe with raw BA 0.98 gave CycleGAN 0.20. Converged probes (validation BA within 0.035 over the last five epochs) agreed (τ 0.82–0.96; per-output ranges at most 0.16, within the widest single-probe interval of 0.18), but every converged probe exceeded the frozen probe's interval on all outputs except HCLD and NeuroCombat: aggressive StarGAN 0.65–0.81 (frozen 0.21), CycleGAN 0.44–0.54 (0.31) and DLEST 0.95–1.00 (0.67–0.71). Orderings changed (τ 0.64–0.73 with the frozen probe): it placed aggressive StarGAN second lowest, converged probes above histogram matching. A probe that has not converged overstates harmonization.

The lowest site accuracy came from destruction. HCLD reached BA 0.07 [0.06, 0.08] under the frozen probe and 0.06–0.11 under every retrained probe, at chance, with the lowest PSNR (13.9 dB) and XCorr (0.71), visible loss of anatomy (Fig. 2), and intensities moved away from the target ($\Delta W_{\text{NYU}} = +0.108$).

Small change is not target alignment. Diffusion translation changed intensities least ($W = 0.008$) yet moved

them barely closer to NYU ($\Delta W_{\text{NYU}} = -0.0013$ [−0.0023, −0.0004], 1% of the raw distance) and raised the KL divergence to NYU from 0.30 to 0.83, whereas histogram matching (−0.053) and CycleGAN (−0.036) moved toward the target. Its outputs are also a random draw: resampling the same checkpoint changed the image more than harmonization did (mean foreground difference 0.085 against 0.075), at unchanged PSNR.

Removed or hidden? Slice probes trained on raw slices reached BA 0.92 on raw test slices (shuffled-label control 0.07 [0.03, 0.10]) but 0.45, 0.33 and 0.41 on histogram-matched, CycleGAN and diffusion outputs, the drop usually reported as harmonization. Probes trained on each method's outputs recovered the source site from its test outputs with BA 0.87, 0.87 and 0.91 (Fig. 1b), with disjoint intervals for every seed. These slice probes use the unconverged benchmark recipe, which understates recoverability (Fig. 1a), so they are conservative bounds. Because head geometry identifies sites, we repeated this on brain-only slices: raw-trained probes reached 0.85 on raw brains but 0.62, 0.48 and 0.39 on the outputs, and probes trained on outputs 0.91, 0.79 and 0.74, again disjoint for every seed with best-validation checkpoints (five of nine with final-epoch ones); brain masks alone gave 0.52–0.54.

Intensity histograms remove geometry. On whole heads, a histogram probe trained on outputs recovered little from histogram matching (0.16 [0.11, 0.22]), whose head histograms are matched by construction, but recovered CycleGAN (0.74) and diffusion (0.70) as well as raw slices (0.74). Restricted to the brain it recovered all three: 0.53 [0.44, 0.63], 0.55 [0.47, 0.64] and 0.38 [0.29, 0.47], against 0.56 on raw brains and 0.13–0.21 for the raw-trained probe. The methods transformed the site signature rather than removing it.

5. DISCUSSION AND RECOMMENDATIONS

Three consequences follow for practice. (i) *One site number is not a verdict.* The same outputs scored 0.13–0.59 under un-converged probes and up to 0.81 under converged ones, and the lowest score belonged to the most destructive output, so site accuracy is readable only beside a probe recipe, evidence of convergence and a fidelity measure. (ii) *Train a probe on the outputs.* Probes trained on harmonized images recovered the source site at 0.87–0.91, and 0.74–0.91 within the brain, where raw-trained probes had suggested removal; what looks like removal is usually concealment. (iii) *Measure the target, not the change.* The output closest to its input moved 1% of the way to the target site. Site labels are themselves confounded: head and brain geometry predict them, so image-probe accuracy cannot certify scanner-effect removal.

We therefore recommend that evaluations (1) report several independently trained probes with their recipe and evidence of convergence, including probes trained on harmonized outputs, and probes restricted to the channel a method should change, with geometry-only controls; (2) report each probe’s raw accuracy and a shuffled-label control; (3) measure alignment to a held-out target reference separately from change, excluding target-site subjects; (4) compute paired uncertainty over subjects, not pixels, and over several draws of stochastic generators; and (5) read site accuracy together with fidelity.

Limitations. We use one cohort, one target site and one axial slice per subject, so volumetric effects are not assessed. Without anatomical ground truth, PSNR and XCorr measure agreement with the input, not correctness. Brain masks remove the head silhouette but not brain shape, slice level or orientation, and site is confounded with demographics and with acquisition geometry that varies within a site; our controls bound these channels but do not separate scanner effects from biological differences. Slice probes use the benchmark recipe, generators were not retuned, and one baseline is transductive.

6. COMPLIANCE WITH ETHICAL STANDARDS

This retrospective study used open-access human data from ABIDE I [10]; ethical approval was not required, as confirmed by the data license.

7. ACKNOWLEDGMENTS

This study used ABIDE I, whose collection was supported by NIMH grants K23MH087770 and R03MH096321, the Stavros Niarchos and Leon Levy Foundations. Experiments used GPU resources of the VPU Lab (UAM). The authors have no relevant financial or non-financial interests to disclose. Generative AI tools (OpenAI Codex, Anthropic Claude) assisted with analysis code and drafting; the authors verified and take full responsibility for all code, results and text.

8. REFERENCES

- [1] J.-P. Fortin et al., “Harmonization of cortical thickness measurements across scanners and sites,” *NeuroImage*, vol. 167, pp. 104–120, 2018.
- [2] B. Glocker et al., “Machine learning with multi-site imaging data: An empirical study on the impact of scanner effects,” *arXiv preprint arXiv:1910.04597*, 2019.
- [3] S. C. Tanaka et al., “A multi-site, multi-disorder resting-state magnetic resonance image database,” *Scientific Data*, vol. 8, pp. 227, 2021.
- [4] B. E. Dewey et al., “DeepHarmony: A deep learning approach to contrast harmonization across scanner changes,” *Magnetic Resonance Imaging*, vol. 64, pp. 160–170, 2019.
- [5] F. Hu et al., “Image harmonization: A review of statistical and deep learning methods for removing batch effects and evaluation metrics for effective harmonization,” *NeuroImage*, vol. 274, pp. 120125, 2023.
- [6] N. K. Dinsdale et al., “Deep learning-based unlearning of dataset bias for MRI harmonisation and confound removal,” *NeuroImage*, vol. 228, pp. 117689, 2021.
- [7] Y. Elazar and Y. Goldberg, “Adversarial removal of demographic attributes from text data,” in *Proc. EMNLP*, 2018, pp. 11–21.
- [8] C. Marzi et al., “Efficacy of MRI data harmonization in the age of machine learning: a multicenter study across 36 datasets,” *Scientific Data*, vol. 11, pp. 115, 2024.
- [9] Saman Rahbar, “Frozen brain-MRI foundation models are site fingerprints,” *arXiv preprint arXiv:2608.10295*, 2026.
- [10] A. Di Martino et al., “The autism brain imaging data exchange: towards a large-scale evaluation of the intrinsic brain architecture in autism,” *Mol. Psychiatry*, vol. 19, no. 6, pp. 659–667, 2014.
- [11] W. E. Johnson et al., “Adjusting batch effects in microarray expression data using empirical Bayes methods,” *Biostat.*, vol. 8, no. 1, pp. 118–127, 2007.
- [12] L. G. Nyúl et al., “New variants of a method of MRI scale standardization,” *IEEE Trans. Med. Imag.*, vol. 19, no. 2, pp. 143–150, 2000.
- [13] J.-Y. Zhu et al., “Unpaired image-to-image translation using cycle-consistent adversarial networks,” in *Proc. ICCV*, 2017, pp. 2223–2232.
- [14] Y. Choi et al., “StarGAN: Unified generative adversarial networks for multi-domain image-to-image translation,” in *Proc. CVPR*, 2018, pp. 8789–8797.
- [15] M. Wu et al., “Disentangled latent energy-based style translation: An image-level structural MRI harmonization framework,” *Neural Netw.*, vol. 184, pp. 107039, 2025.
- [16] J. Ho et al., “Denoising diffusion probabilistic models,” in *Proc. NeurIPS*, 2020, vol. 33, pp. 6840–6851.
- [17] C. Meng et al., “SDEdit: Guided image synthesis and editing with stochastic differential equations,” in *Proc. ICLR*, 2022.
- [18] M. Wu et al., “Unpaired volumetric harmonization of brain MRI with conditional latent diffusion,” *Med. Image Anal.*, vol. 107, pp. 103849, 2026.
- [19] K. He et al., “Deep residual learning for image recognition,” in *Proc. CVPR*, 2016, pp. 770–778.
- [20] F. Isensee et al., “Automated brain extraction of multisequence MRI using artificial neural networks,” *Hum. Brain Mapp.*, vol. 40, no. 17, pp. 4952–4964, 2019.