

FROZEN SITE-PROBE ACCURACY IS NOT A STANDALONE HARMONIZATION SCORE: AN AUDIT OF EVALUATION PRACTICE FOR MULTI-SITE MRI

Muhammad Muqsit Islam^{1,2,*†} Gloria García Cuenco^{1,2,†}
José M. Martínez Sánchez¹ Pierrick Coupé³

¹ Universidad Autónoma de Madrid, Spain ² University of Bordeaux, France

³ Université de Bordeaux, CNRS, Bordeaux INP, LaBRI, UMR 5800, F-33400 Talence, France

ABSTRACT

Image-level MRI harmonization is often judged with a site probe, a classifier that predicts the acquisition site: lower accuracy on harmonized images is read as success. We examine how reliable such verdicts are, starting from a frozen probe trained on unharmonized images, on raw ABIDE I T1-weighted MRI (1,112 subjects, 17 sites) with ten harmonized outputs as test cases, one from a published model (HACA3) with its pretrained weights. Verdicts depended on probe training: five seeds of one recipe gave the same CycleGAN output balanced accuracies from 0.13 to 0.59, and five longer-trained, stable probes found the site more decodable than the frozen probe in eight of ten outputs. A Gaussian blur, moving intensities away from the target, scored lowest on the frozen probe, at the shuffled-label baseline. Probes trained on harmonized images decoded the source site far better than raw-trained probes (+0.42 to +0.55 balanced accuracy, paired intervals above zero), also for HACA3, which held raw-trained probes near 0.28, and on brain-only slices and their intensity histograms, so low frozen-probe accuracy does not establish removal of site-label information. The output with the smallest intensity-distribution change barely moved toward the target. We recommend several well-trained probes, including output-trained ones, against a shuffled-label baseline, with target alignment reported apart from image change.

Index Terms— MRI harmonization, evaluation, site classification, domain shift, reproducibility

1. INTRODUCTION

Acquisition differences between sites confound pooled MRI analyses and degrade models on unseen scanners [1, 2]. Image-level harmonization translates scans toward a common appearance. Travelling subjects [3] or paired protocols [4] allow direct checks on few scanners, so large retrospective cohorts are judged indirectly [5]. A common proxy is a *site probe*: a classifier that predicts the acquisition site from

an image, which is easy on raw MRI [2]. If its accuracy drops on harmonized images, site effects are taken to be removed [6, 7]; this is reported alongside similarity to the input, such as PSNR, and intensity-distribution distances. Some studies retrain the probe on harmonized images [6] or on image features [7, 8]. Others train it on original images and apply it unchanged to harmonized ones: the PRISM method [9], for example, reports this accuracy drop, with anatomy and segmentation checks, as evidence of harmonization. This frozen-probe protocol is our starting point.

Each proxy measures less than it appears to. A frozen probe reports what one classifier recovers; it can also fail because outputs are unfamiliar, damaged or noisy. Removing an attribute until a fixed adversary fails does not stop a new classifier from recovering it [10]. Site is decodable at about 0.9 chance-corrected balanced accuracy even from randomly initialized encoders [11]. Similarity to the input rewards doing nothing, and small change says nothing about the target.

We examine how reliable site-probe verdicts are, starting from the frozen probe of an earlier benchmark, on outputs of seven harmonization families, one a published model as released. Site classifiers retrained after ComBat [12] or histogram matching [13] still predict site; we quantify how probe training changes the verdict and whether low frozen-probe accuracy survives probes trained on harmonized outputs, separating geometry from intensity, measuring target alignment and adding a blur control. We contribute a protocol with paired uncertainty and probe controls, these findings and a reporting checklist (Sec. 5).¹ We neither propose nor rank harmonizers.

2. DATA AND AUDITED OUTPUTS

Cohort. Raw ABIDE I T1-weighted scans [14] of 1,112 subjects from 17 sites are split by subject, stratified by site, into 890 training, 113 validation and 109 test subjects. Volumes are reoriented (in-plane rotation still varies), clipped to $[p_1, p_{99}]$ and scaled to $[0, 1]$; one axial slice per subject is

*Corresponding author

† Equal contribution

¹Code, the evaluation protocol with all amendments, and per-subject results: <https://github.com/muqsitamir/HarmonIt>.

chosen by foreground area, cropped to a dilated head mask and resized to 256×256 . NYU, the largest site, is the target; 19 test subjects are from NYU and 90 from 16 source sites. Slices are frozen and verified for every output. Labels are coarser than acquisitions: UM, UCLA and Leuven each merge two released sub-samples, and 11 of 17 sites mix voxel sizes or matrices; no per-subject scanner identifier is released.

Audited outputs. Earlier benchmark outputs were evaluated without retraining: NeuroCombat [1, 15] per pixel, fit on the test cohort itself (transductive); histogram matching of foreground quantiles to a reference pooled over all 17 training sites, not NYU, as in the benchmark [16]; CycleGAN [17]; StarGAN [18], aggressive and conservative; a DLEST-style baseline without DLEST’s energy-based sampler [6], after 1,000 and 1,500 steps; conditional diffusion image-to-image translation [19, 20] (strength 0.35, 50 DDIM steps), sampled twice from one checkpoint. We add HACA3 [21] as released (pretrained weights, NYU target image) after its required N4 correction and rigid MNI registration, mapped back to native space: ten outputs (Table 1 omits DLEST-style 1500 and conservative StarGAN). As a control that ignores the target, raw test slices were blurred with a Gaussian ($\sigma = 2, 4$ and 8 pixels). We trained the other learned methods on the training split. Test subjects were inspected in the earlier benchmark, so this is a retrospective reanalysis; each audit experiment was written into a protocol before its outcome was seen.

3. AUDIT PROTOCOL

Change and target alignment. PSNR on a fixed $[0, 1]$ scale and normalized cross-correlation (XCorr) compare output with input on whole images; most target-conditioned generators return NYU inputs almost unchanged, so we summarize the 90 source subjects. *Change* is the per-subject Wasserstein-1 distance W between input and output intensities. *Target alignment* is the Wasserstein-1 distance and KL divergence between foreground intensities (defined on the raw slice) and a reference pooled from all 147 NYU training subjects; ΔW_{NYU} is the output’s distance minus the input’s.

Site probes. Image probes are single-channel ResNet-18s [22] predicting 17 sites, tested on the fixed slices; unless stated otherwise they use AdamW (learning rate 3×10^{-4} , batch 64, 500 steps) and random affine augmentation and are selected on validation BA. We use 53 probes in five families, 44 of them ResNets trained by us, and 16 for HACA3:

Frozen (1): the earlier benchmark’s own checkpoint.

Benchmark recipe (5 seeds): that training repeated, on a random axial slice of a raw training volume at each step. Does the seed change the verdict?

Converged (5 seeds): four times the steps, cosine decay, reseeded data-loader workers, final epoch. Do they agree?

Shuffled labels (4): labels permuted between subjects. Is the probe reading images, not labels?




Fig. 1. Probes trained on raw slices (grey) or a method’s outputs (aqua), tested on its outputs (diffusion: draw 2), on (a) whole heads and (b) brains: image probes (circles, mean and range over 3 seeds) and intensity-histogram probes (squares, 95% interval). Numbers: paired output minus raw-trained difference. Black bar: raw-trained probe on raw slices. Grey: silhouette or brain-mask probes.

Slice probes (24, plus 6 restricted-input networks and 8 logistic fits): trained on one fixed slice per training subject, from raw images or one method’s outputs. Is site still decodable? Restricted variants see only head silhouettes or brain masks (geometry), only foreground intensity histograms (intensity; logistic regression), or slices inside raw-volume HD-BET brain masks [23].

Slice-probe methods were set in advance (intensity, adversarial, diffusion); converged, silhouette, histogram and brain-only probes came after earlier results but before their own.

Statistics. Source BA averages the 16 source sites (2–15 test subjects each). A uniform guess over the 17 classes scores $1/17 \approx 0.06$ and one source class predicted for every subject $1/16$; shuffled-label probes, our empirical null, scored 0.03–0.08. Intervals are 95% percentiles of 2,000 within-site subject bootstrap replicates, paired across methods and probes, and conditional on these sites and checkpoints.

4. RESULTS

Probe validity. Raw source BA was 0.95 for the frozen probe, 0.81–0.98 for benchmark-recipe and 1.00 for converged probes. Probes learn more than scanner appearance: head silhouettes alone gave 0.73–0.86, and age, sex and diagnosis 0.22 [0.15, 0.29] (0.31 with brain volume).

Table 1. Source-subject means ($n = 90$). PSNR, XCorr, W : output vs. input; ΔW_{NYU} (negative is closer), KL_{NYU} : distance to NYU. Site BA: frozen probe with 95% interval and share labelled NYU; 5 benchmark-recipe and 5 converged probes: mean [min, max]; blur: frozen probe only. Shuffled labels: 0.03–0.08. [†]Fit on the test cohort. [‡]Pooled 17-site reference, not NYU.

Method	Change			Target alignment		Frozen probe		Probes retrained on raw images	
	PSNR	XCorr	W	ΔW_{NYU}	KL_{NYU}	Source BA	NYU	Benchmark recipe	Converged recipe
Raw input	–	–	0	0	0.30	0.95	0%	0.90 [0.81, 0.98]	1.00 [1.00, 1.00]
NeuroCombat [†]	20.0	0.940	0.055	+0.003	0.88	0.27 [0.21, 0.32]	2%	0.33 [0.26, 0.37]	0.34 [0.28, 0.42]
Histogram matching [‡]	23.0	0.982	0.054	-0.053	0.09	0.44 [0.37, 0.52]	10%	0.44 [0.27, 0.59]	0.61 [0.55, 0.68]
CycleGAN	21.3	0.977	0.048	-0.036	0.25	0.31 [0.25, 0.37]	52%	0.28 [0.13, 0.59]	0.49 [0.44, 0.54]
StarGAN (aggr.)	16.0	0.896	0.071	-0.019	0.32	0.21 [0.15, 0.27]	56%	0.35 [0.18, 0.54]	0.70 [0.65, 0.81]
DLEST-style 1000	28.4	0.993	0.017	-0.009	0.31	0.71 [0.65, 0.77]	6%	0.84 [0.76, 0.89]	0.99 [0.96, 1.00]
Diffusion, draw 1	21.7	0.948	0.008	-0.001	0.83	0.37 [0.32, 0.42]	36%	0.42 [0.24, 0.63]	0.57 [0.52, 0.59]
Diffusion, draw 2	21.7	0.948	0.008	-0.002	0.82	0.35 [0.29, 0.40]	38%	0.42 [0.30, 0.62]	0.56 [0.50, 0.59]
HACA3 (pretrained)	18.9	0.954	0.062	-0.016	2.29	0.25 [0.18, 0.32]	40%	0.29 [0.18, 0.37]	0.28 [0.23, 0.33]
Blur control, $\sigma=2$	24.7	0.974	0.018	+0.016	0.57	0.07 [0.06, 0.10]	0%	–	–




Fig. 2. Median-PSNR source subject and an axial NYU slice with similar head outline and anatomy, in the input’s window; insets: 2× zoom of the box. Bottom: output – input with PSNR. Diffusion texture differs by draw; blur has the highest PSNR.

Verdicts depend on probe training. Benchmark-recipe probes assigned very different site BA to the same outputs (Table 1): CycleGAN 0.13–0.59 and aggressive StarGAN 0.18–0.54, with Kendall τ 0.39–0.83 between method orderings. Raw accuracy did not reveal this: the probe with raw BA 0.98 gave CycleGAN 0.20. Much of this spread was training instability: converged probes (validation BA within 0.035 over the last five epochs) varied far less (τ 0.78–0.94; per-output ranges at most 0.16), but every converged probe exceeded the frozen probe’s interval on all outputs except NeuroCombat and HACA3, e.g. aggressive StarGAN 0.65–0.81 against 0.21. Orderings changed (τ 0.56–0.67 with the frozen probe).

A blur scored lower than every harmonizer. A $\sigma = 2$ blur kept PSNR at 24.7 dB, above CycleGAN and diffusion, and moved intensities away from NYU ($\Delta W_{\text{NYU}} = +0.016$), yet scored frozen-probe BA 0.07 [0.06, 0.10], in the shuffled-label range ($\sigma = 4, 8$: 0.06, 0.07). It labelled 87 of 90 blurred slices Yale, none NYU; on CycleGAN, aggressive StarGAN and diffusion 65–80% of errors named NYU. A perturbation ignoring the target can collapse frozen-probe accuracy.

Small change is not target alignment. Diffusion changed the intensity distribution least ($W = 0.008$) yet barely moved it toward NYU ($\Delta W_{\text{NYU}} = -0.0013$

[−0.0023, −0.0004], 1% of the raw distance) and raised KL divergence to NYU from 0.30 to 0.83. A second sample from the same checkpoint changed images more than harmonization did (median absolute foreground difference 0.085 vs. 0.075), at unchanged PSNR.

Is site still decodable? Slice probes trained on raw slices reached BA 0.92 on raw test slices but 0.45, 0.33 and 0.41 on histogram-matched, CycleGAN and diffusion outputs, a drop interpreted as evidence of harmonization [9]. Probes trained on each method’s outputs recovered the source site from its test outputs with BA 0.87, 0.87 and 0.91 (Fig. 1a), exceeding raw-trained probes by +0.42 [0.36, 0.48], +0.55 [0.52, 0.58] and +0.50 [0.45, 0.54] (paired differences of unrounded three-seed means; benchmark recipe, so a lower bound). Because head geometry identifies sites, we repeated this on brain-only slices: raw-trained probes reached 0.85 on raw brains but 0.62, 0.48 and 0.39 on the outputs, and probes trained on outputs 0.91, 0.79 and 0.74 (Fig. 1b; differences +0.29 to +0.36, all intervals above zero, also with final-epoch checkpoints); brain masks alone gave 0.52–0.54. From intensity histograms alone, output-trained logistic probes recovered the site from CycleGAN and diffusion heads (0.74, 0.70; raw 0.74) and, within the brain, from all three outputs (0.38–0.55; raw brains 0.56; raw-trained probe 0.13–0.21), but little from histogram-

matched heads (0.16), matched by construction. Low frozen-probe accuracy does not establish removal of site-label information, much of which output-trained probes decode.

A published harmonizer. HACA3 kept every raw-trained probe low (frozen 0.25 [0.18, 0.32], benchmark recipe 0.18–0.37, converged 0.23–0.33, slice probes 0.28; 54% of frozen-probe errors named NYU), yet probes trained on its outputs recovered the source site at 0.75 (+0.47 [0.43, 0.52]). Its raw-trained probes were retrained with the original recipes and seeds: converged ones reproduced the originals (median per-output difference 0.02), benchmark-recipe ones did not (up to 0.32). Its KL reflects CSF rendered near zero.

5. DISCUSSION AND RECOMMENDATIONS

One site number is not a verdict: it moved by up to 0.46 across seeds and was lowest for a blur. We recommend reporting (1) several converged probes, their recipe, raw accuracy, a shuffled-label null and a fidelity measure; (2) probes trained on the outputs, with channel-restricted and geometry controls, whose high accuracy shows retained site-label information, not by itself failed harmonization, as head and brain shape predict site; (3) intensity alignment to the target apart from change, on source subjects: the output with the smallest change closed 1% of its distance to NYU; and (4) paired uncertainty over subjects and generator draws.

Limitations. We use one cohort, one target site and one axial slice per subject (no volumetric effects), with 2–15 test subjects per source site. Without anatomical ground truth, PSNR and XCorr measure agreement with the input, not correctness. Site is confounded with demographics, brain shape, slice level and acquisition geometry; our controls bound these channels but cannot separate scanner from biological effects. Generators were not retuned; one baseline is transductive.

6. COMPLIANCE WITH ETHICAL STANDARDS

This research study was conducted retrospectively using human subject data made available in open access by ABIDE I [14], whose sites shared the fully anonymized data with local ethics approval or an explicit waiver. Ethical approval was not required as confirmed by the license attached with the open access data.

7. ACKNOWLEDGMENTS

ABIDE I was supported by NIMH grants K23MH087770 and R03MH096321 and the Stavros Niarchos and Leon Levy Foundations. This work also benefited from the support of the project HoliBrain of the French National Research Agency (ANR-23-CE45-0020-01). Finally, this study received financial support from the French government in the framework of the PEPR StratifyAging, the PEPR Prodrom-ND and the IHU VBHI (ANR-23-IAHU-0001). The authors declare no

conflicts of interest. OpenAI Codex and Anthropic Claude helped write code (Secs. 3–4, Table 1, Figs. 1–2) and draft the text; the authors designed the study, reviewed the code, verified all numbers and take full responsibility.

8. REFERENCES

- [1] J.-P. Fortin et al., “Harmonization of cortical thickness measurements across scanners and sites,” *NeuroImage*, vol. 167, pp. 104–120, 2018.
- [2] B. Glocker et al., “Machine learning with multi-site imaging data: An empirical study on the impact of scanner effects,” *arXiv preprint arXiv:1910.04597*, 2019.
- [3] S. C. Tanaka et al., “A multi-site, multi-disorder resting-state magnetic resonance image database,” *Scientific Data*, vol. 8, pp. 227, 2021.
- [4] B. E. Dewey et al., “DeepHarmony: A deep learning approach to contrast harmonization across scanner changes,” *Magnetic Resonance Imaging*, vol. 64, pp. 160–170, 2019.
- [5] F. Hu et al., “Image harmonization: A review of statistical and deep learning methods for removing batch effects and evaluation metrics for effective harmonization,” *NeuroImage*, vol. 274, pp. 120125, 2023.
- [6] M. Wu et al., “Disentangled latent energy-based style translation: An image-level structural MRI harmonization framework,” *Neural Netw.*, vol. 184, pp. 107039, 2025.
- [7] M. Wu et al., “Unpaired volumetric harmonization of brain MRI with conditional latent diffusion,” *Med. Image Anal.*, vol. 107, pp. 103849, 2026.
- [8] N. K. Dinsdale et al., “Deep learning-based unlearning of dataset bias for MRI harmonisation and confound removal,” *NeuroImage*, vol. 228, pp. 117689, 2021.
- [9] S. Galada et al., “PRISM: Privacy-preserving inter-site MRI harmonization via disentangled representation learning,” in *Proc. ISBI*, 2025.
- [10] Y. Elazar and Y. Goldberg, “Adversarial removal of demographic attributes from text data,” in *Proc. EMNLP*, 2018, pp. 11–21.
- [11] S. Rahbar, “Frozen brain-MRI foundation models are site fingerprints,” *arXiv preprint arXiv:2608.10295*, 2026.
- [12] C. Marzi et al., “Efficacy of MRI data harmonization in the age of machine learning: a multicenter study across 36 datasets,” *Scientific Data*, vol. 11, pp. 115, 2024.
- [13] R. Souza et al., “Image-encoded biological and non-biological variables may be used as shortcuts in deep learning models trained on multi-site neuroimaging data,” *J. Am. Med. Inform. Assoc.*, vol. 30, no. 12, pp. 1925–1933, 2023.
- [14] A. Di Martino et al., “The autism brain imaging data exchange: towards a large-scale evaluation of the intrinsic brain architecture in autism,” *Mol. Psychiatry*, vol. 19, no. 6, pp. 659–667, 2014.
- [15] W. E. Johnson et al., “Adjusting batch effects in microarray expression data using empirical Bayes methods,” *Biostat.*, vol. 8, no. 1, pp. 118–127, 2007.
- [16] L. G. Nyúl et al., “New variants of a method of MRI scale standardization,” *IEEE Trans. Med. Imag.*, vol. 19, no. 2, pp. 143–150, 2000.
- [17] J.-Y. Zhu et al., “Unpaired image-to-image translation using cycle-consistent adversarial networks,” in *Proc. ICCV*, 2017, pp. 2223–2232.
- [18] Y. Choi et al., “StarGAN: Unified generative adversarial networks for multi-domain image-to-image translation,” in *Proc. CVPR*, 2018, pp. 8789–8797.
- [19] J. Ho et al., “Denoising diffusion probabilistic models,” in *Proc. NeurIPS*, 2020, vol. 33, pp. 6840–6851.
- [20] C. Meng et al., “SDEdit: Guided image synthesis and editing with stochastic differential equations,” in *Proc. ICLR*, 2022.
- [21] L. Zuo et al., “HACA3: A unified approach for multi-site MR image harmonization,” *Comput. Med. Imaging Graph.*, vol. 109, pp. 102285, 2023.
- [22] K. He et al., “Deep residual learning for image recognition,” in *Proc. CVPR*, 2016, pp. 770–778.
- [23] F. Isensee et al., “Automated brain extraction of multisequence MRI using artificial neural networks,” *Hum. Brain Mapp.*, vol. 40, no. 17, pp. 4952–4964, 2019.